

Bayesian Inverse Problems in PDEs with Probabilistic Models for Numerical Error

Jon Cockayne¹

Chris Oates²

Tim Sullivan³

Mark Girolami⁴

¹ University of Warwick

² University of Technology Sydney

³ Free University of Berlin

⁴ Alan Turing Institute for Data Science

February 2, 2016

What is an “inverse problem”?

$$y = \mathcal{G}(\theta) + \xi$$

- $\theta \in \mathbb{R}^n$ parameters
- \mathcal{G} parameter to observable map
- $y \in \mathbb{R}^m$ observations

Inverse Problem: Given observations y , what is the ‘value’ of θ ?

Bayesian inverse problem. Given observations y **and a prior for θ** , **what is the posterior?**

What is a “Probabilistic model for numerical error”?

Roots in Diaconis [1988].

Recent interest in:

- Quadrature (Briol et al. [2015])
- Optimization (Snoek et al. [2012])
- ODEs and PDEs (Conrad et al. [2015], Schober et al. [2014])

What is a “Probabilistic model for numerical error”?

Roots in Diaconis [1988].

Recent interest in:

- Quadrature (Briol et al. [2015])
- Optimization (Snoek et al. [2012])
- ODEs and PDEs (Conrad et al. [2015], Schober et al. [2014])

ProbNum attempts to give **probabilistic answers** to ‘**deterministic**’ **problems**.

Deterministic method \nrightarrow accurate solution.

1 Motivation

2 Linear Problems

Example: Canonical Elliptic PDE

3 Nonlinear Problems

Example: Steady-State Allen-Cahn Equation

4 Summary

Motivation

Our \mathcal{G} is of the form:

$$\begin{aligned}\mathcal{A}u(\mathbf{x}) &= g(\mathbf{x}) & \mathbf{x} \in D \\ \mathcal{B}u(\mathbf{x}) &= b(\mathbf{x}) & \mathbf{x} \in \partial D\end{aligned}$$

\mathcal{A} , \mathcal{B} are **differential operators**. Think of this as a system of operator equations:

$$\mathcal{O}u := \begin{bmatrix} \mathcal{A} \\ \mathcal{B} \end{bmatrix} u = \begin{bmatrix} g \\ b \end{bmatrix}$$

Example: Darcy Flow

$$\mathcal{A}u := \nabla \cdot \kappa(\mathbf{x}) \nabla u(\mathbf{x})$$

$$\mathcal{B}u := u$$

Why ProbNum?

Finite Element method constructs the solution over elements of a mesh.

“Gridding away” discretisation error is often computationally infeasible:

- Complex, high-dimensional domain.
- Complex, nonlinear, parabolic PDE.

Instead try to **capture** it.

Why ProbNum?

Finite Element method constructs the solution over elements of a mesh.

“Gridding away” discretisation error is often computationally infeasible:

- Complex, high-dimensional domain.
- Complex, nonlinear, parabolic PDE.

Instead try to **capture** it.

However. . . probabilistic approach may make already costly methods more costly. Why bother?

- Inverse Problems:
 - Incorporate covariance into likelihood.
 - Posterior distribution reflects numerical error.

Linear Problems

$$\mathcal{A}u(\mathbf{x}) = g(\mathbf{x}) \quad \mathbf{x} \in D$$

$$\mathcal{B}u(\mathbf{x}) = b(\mathbf{x}) \quad \mathbf{x} \in \partial D$$

- \mathcal{A}, \mathcal{B} linear (for now). We build a Gaussian Process model for u .
- Prior: $u \sim GP(0, k)$ for some kernel k .

$$\mathcal{A}u(\mathbf{x}) = g(\mathbf{x}) \quad \mathbf{x} \in D$$

$$\mathcal{B}u(\mathbf{x}) = b(\mathbf{x}) \quad \mathbf{x} \in \partial D$$

- \mathcal{A}, \mathcal{B} linear (for now). We build a Gaussian Process model for u .
- Prior: $u \sim GP(0, k)$ for some kernel k .
- For subsets $X = \{x_i\}, Y = \{y_j\}$ of \bar{D} :

$$K(X, Y) = [k(x_i, y_j)]_{ij}$$

$$\mathcal{A}K(X, Y) = [\mathcal{A}k(x_i, y_j)]_{ij} \quad \text{etc.}$$

$$\mathcal{A}u(\mathbf{x}) = g(\mathbf{x}) \quad \mathbf{x} \in D$$

$$\mathcal{B}u(\mathbf{x}) = b(\mathbf{x}) \quad \mathbf{x} \in \partial D$$

- \mathcal{A}, \mathcal{B} linear (for now). We build a Gaussian Process model for u .
- Prior: $u \sim GP(0, k)$ for some kernel k .
- For subsets $X = \{x_i\}$, $Y = \{y_j\}$ of \bar{D} :

$$K(X, Y) = [k(x_i, y_j)]_{ij}$$

$$\mathcal{A}K(X, Y) = [\mathcal{A}k(x_i, y_j)]_{ij} \quad \text{etc.}$$

- $\bar{\mathcal{A}}$ denotes the adjoint of \mathcal{A} :

$$\mathcal{A}k(x, y) \quad \bar{\mathcal{A}}k(x, y)$$

Choose design points $X_0 = X_0^A \cup X_0^B$ & observations $\mathbf{v} = \{v_i\} \dots$

Choose design points $X_0 = X_0^{\mathcal{A}} \cup X_0^{\mathcal{B}}$ & observations $\mathbf{v} = \{v_i\} \dots$

Posterior Measure for u

$$u(\mathbf{x})|\mathbf{v} \sim \mathcal{N}(\mu, \Sigma)$$

where:

$$\mu := \bar{O}K(\mathbf{x}, X_0) [\bar{O}\bar{O}K(X_0, X_0)]^{-1} \mathbf{v}$$

$$\Sigma := K(\mathbf{x}, \mathbf{x}) - \bar{O}K(\mathbf{x}, X_0) [\bar{O}\bar{O}K(X_0, X_0)]^{-1} \bar{O}K(X_0, \mathbf{x})$$

Choose design points $X_0 = X_0^A \cup X_0^B$ & observations $\mathbf{v} = \{v_i\} \dots$

Posterior Measure for u

$$u(\mathbf{x})|\mathbf{v} \sim \mathcal{N}(\mu, \Sigma)$$

where:

$$\mu := \bar{O}K(\mathbf{x}, X_0) [\mathcal{O}\bar{O}K(X_0, X_0)]^{-1} \mathbf{v}$$

$$\Sigma := K(\mathbf{x}, \mathbf{x}) - \bar{O}K(\mathbf{x}, X_0) [\mathcal{O}\bar{O}K(X_0, X_0)]^{-1} \mathcal{O}K(X_0, \mathbf{x})$$

$$\mathcal{O}K(X_0, \mathbf{x}) = \begin{bmatrix} \mathcal{A}K(X_0^A, \mathbf{x}) \\ \mathcal{B}K(X_0^B, \mathbf{x}) \end{bmatrix}$$

$$\mathcal{O}\bar{O}K(X, Y) = \begin{bmatrix} \mathcal{A}\bar{\mathcal{A}}K(X_0^A, X_0^A) & \mathcal{A}\bar{\mathcal{B}}K(X_0^A, X_0^B) \\ \mathcal{B}\bar{\mathcal{A}}K(X_0^B, X_0^A) & \mathcal{B}\bar{\mathcal{B}}K(X_0^B, X_0^B) \end{bmatrix}$$

Linear Case: Inverse Problem

We want to use a likelihood like:

$$p(\mathbf{y}|\theta) = \prod p(y_i|u(x_i; \theta))$$

In practise we usually use:

$$\hat{p}(\mathbf{y}|\theta) = \prod p(y_i|\hat{u}(x_i; \theta))$$

Linear Case: Inverse Problem

We want to use a likelihood like:

$$p(\mathbf{y}|\theta) = \prod p(y_i|u(x_i; \theta))$$

In practise we usually use:

$$\hat{p}(\mathbf{y}|\theta) = \prod p(y_i|\hat{u}(x_i; \theta))$$

Instead we will marginalise u :

$$\begin{aligned} p_{PN}(\mathbf{y}|\theta) &= \int \prod p(y_i|u(x_i; \theta))p(u)du \\ &= N(\mu, \sigma^2 I + \Sigma) \end{aligned}$$


Linear Example

$$\begin{aligned}\nabla \cdot \kappa(\mathbf{x}; \theta) \nabla u(\mathbf{x}) &= 0 && \text{in } [0, 1]^2 \\ u(\mathbf{x}) &= x_1 && \text{at } x_2 = 0 \\ u(\mathbf{x}) &= 1 - x_1 && \text{at } x_2 = 1 \\ \frac{\partial u}{\partial x_1} &= 0 && \text{at } x_1 = 0, 1\end{aligned}$$

Approximate κ with a truncated Karhunen-Loève expansion:

$$\kappa(\mathbf{x}; \theta) = \exp \left(\sum_{k=1}^6 \frac{\theta_k}{k^2} \cos(2\pi(s_k x_1 + r_k x_2)) \right)$$

Canonical Elliptic PDE: Solution




Left: Model solution (generated with FEM, 100×100 mesh).

Right: Mean function of PN solution, design points annotated.


PN solution uses a **squared exponential** kernel: $k = \exp\left(-\frac{\|x-y\|_2^2}{2\sigma^2}\right)$

Forward Sample Paths



Samples drawn from posterior on a 10×10 grid, along $x_1 = x_2$.

Inverse Problem



Posterior distributions of coefficients of κ , using the PN forward model (left) vs. FEM (right)

Dashed lines / black crosses reflect the **true values**.

Nonlinear Problems

$$\mathcal{A}u(\mathbf{x}) = g(\mathbf{x}) \quad \mathbf{x} \in D$$

$$\mathcal{B}u(\mathbf{x}) = b(\mathbf{x}) \quad \mathbf{x} \in \partial D$$

- \mathcal{A} and/or \mathcal{B} **nonlinear**. No more nice closed form posterior!
- Think about a class of PDEs which decompose as:

$$(\mathcal{L} + \mathcal{A})u = g$$

\mathcal{L} linear, \mathcal{A} **invertible**.

Nonlinear Case: Theory II

- Introduce a 'latent function' z
- Construct a **new system**:

$$\begin{array}{ccccc} (\mathcal{L} + \mathcal{A})u = g & & \mathcal{L}u = (g - z) & & \mathcal{L}u = (g - z) \\ & \longrightarrow & \mathcal{A}u = z & \longrightarrow & u = \mathcal{A}^{-1}z \\ \mathcal{B}u = b & & \mathcal{B}u = b & & \mathcal{B}u = b \end{array}$$

Nonlinear Case: Theory II

- Introduce a 'latent function' z
- Construct a **new system**:

$$\begin{array}{ccccc} (\mathcal{L} + \mathcal{A})u = g & & \mathcal{L}u = (g - z) & & \mathcal{L}u = (g - z) \\ \mathcal{B}u = b & \longrightarrow & \mathcal{A}u = z & \longrightarrow & u = \mathcal{A}^{-1}z \\ & & \mathcal{B}u = b & & \mathcal{B}u = b \end{array}$$

- Integrate out z :
 - Generate an **approximate solution** \hat{u} (eg. using FEM on a coarse mesh)
 - Guess $z = \mathcal{A}\hat{u}$.
 - Place an **importance distribution** over z .

Nonlinear Case: Inverse Problem

Can no longer evaluate p_{PN} exactly. Need to marginalize z :

$$p_{PN}(y|\theta) = \int \frac{1}{p(z|\theta)} \int p(y|u(x, \theta)) p(u) du p(z|\theta) dz$$

The outer integral is not exactly solvable - use an unbiased estimate.

\implies pseudo-marginal MCMC.




Steady-State Allen-Cahn Equation

$$\begin{aligned} -\delta \nabla^2 u + \delta^{-1}(u^3 - u) &= 0 & (x, y) \in (0, 1)^2 \\ u &= 1 & x = 0, x = 1, y \in (0, 1) \\ u &= -1 & y = 0, y = 1 \end{aligned}$$

Steady-State Allen-Cahn Equation

$$\begin{aligned} -\delta \nabla^2 u + \delta^{-1}(u^3 - u) &= 0 & (x, y) \in (0, 1)^2 \\ u &= 1 & x = 0, x = 1, y \in (0, 1) \\ u &= -1 & y = 0, y = 1 \end{aligned}$$

Three solutions!



Inverse Problem

We want to infer δ .

Multiple solutions makes the inverse problem **very hard!**

However...in the PN setting we can place a **mixture** distribution over z .

Inverse Problem


We want to infer δ .

Multiple solutions makes the inverse problem **very hard!**

However... in the PN setting we can place a **mixture** distribution over z .

Right: Inference of δ using PN.

- Prior: $\delta \sim U(0.02, 1)$.
- Deflation [Farrell et al., 2014] used to generate the multiple \hat{u} at each δ .



Summary

- With PN we can capture and propagate uncertainty from numerical error.
- This helps us to **make sound inferences!**
- Many sampling challenges with nonlinear problems.

Lots still to be done. . .

- Computation — very expensive compared to FEM
- Development needed on the weak form [Conrad et al., 2015].
- Optimal choices of kernel / RKHS [Owhadi, 2014].
- Experimental design — how should we place X_0 to minimise posterior uncertainty?
- Applications to physical problems of interest.

Thanks!

References

- François-Xavier Briol, Chris J Oates, Mark Girolami, Michael A Osborne, and Dino Sejdinovic. Probabilistic integration. *arXiv preprint arXiv:1512.00933*, 2015.
- Patrick R Conrad, Mark Girolami, Simo Särkkä, Andrew Stuart, and Konstantinos Zygalakis. Probability measures for numerical solutions of differential equations. *arXiv preprint arXiv:1506.04592*, 2015.
- Persi Diaconis. Bayesian numerical analysis. *Statistical decision theory and related topics IV*, 1:163–175, 1988.
- Patrick E Farrell, Ásgeir Birkisson, and Simon W Funke. Deflation techniques for finding distinct solutions of nonlinear partial differential equations. *arXiv preprint arXiv:1410.5620*, 2014.
- Houman Owhadi. Bayesian numerical homogenization. *arXiv preprint arXiv:1406.6668*, 2014.
- Michael Schober, David K Duvenaud, and Philipp Hennig. Probabilistic ode solvers with runge-kutta means. In *Advances in Neural Information Processing Systems*, pages 739–747, 2014.
- Jasper Snoek, Hugo Larochelle, and Rp Adams. Practical Bayesian Optimization of Machine Learning Algorithms. *Nips*, pages 1–9, 2012. doi: 2012arXiv1206.2944S. URL <https://papers.nips.cc/paper/4522-practical-bayesian-optimization-of-machine-learning-algorithms.pdf>.